

2348441_lab_05

March 21, 2024

Lab Exercise 5 - Data Exploration Parametric Methods

- Created by : Nileem Kaveramma C C | 2348441
- Created DATE:20-03-2024
- Edited Date: 21-03-2024

EMPLOYEE SALARY ANALYSIS The provided dataset captures information relevant to employee salary prediction, encompassing various attributes such as age, gender, education level, job title, years of experience, and salary. With a diverse set of features, the dataset offers valuable insights into the characteristics of individuals within an organizational context. This dataset becomes particularly relevant for exploring patterns and relationships that could contribute to predicting employee salaries. Through descriptive statistics, visualizations, and parametric tests, analysts can discern trends, potential disparities, and factors influencing salary variations among employees.

IMPORTED LIBRARIES

- numpy - for numerical, array, matrices (Linear Algebra) processing
- Pandas - for loading and processing datasets
- matplotlib.pyplot - For visualisation
- Saeborn - for statistical graph

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

df is a commonly used variable name that often represents a DataFrame

```
[ ]: df = pd.read_csv('/content/Salary Data.csv')
df
```

```
[ ]:      Age  Gender Education Level      Job Title \
0    32.0   Male   Bachelor's      Software Engineer
1    28.0  Female   Master's      Data Analyst
2    45.0   Male      PhD      Senior Manager
3    36.0  Female   Bachelor's      Sales Associate
4    52.0   Male   Master's      Director
..    ...    ...    ...    ...
370  35.0  Female   Bachelor's      Senior Marketing Analyst
371  43.0   Male   Master's      Director of Operations
```

372	29.0	Female	Bachelor's	Junior Project Manager
373	34.0	Male	Bachelor's	Senior Operations Coordinator
374	44.0	Female	PhD	Senior Business Analyst

	Years of Experience	Salary
0	5.0	90000.0
1	3.0	65000.0
2	15.0	150000.0
3	7.0	60000.0
4	20.0	200000.0
..
370	8.0	85000.0
371	19.0	170000.0
372	2.0	40000.0
373	7.0	90000.0
374	15.0	150000.0

[375 rows x 6 columns]

df.shape - attribute is used to get the dimensions of the DataFrame.

```
[ ]: df.shape
```

```
[ ]: (375, 6)
```

df.columns attribute is used to retrieve the column labels or names of the DataFrame

```
[ ]: df.columns
```

```
[ ]: Index(['Age', 'Gender', 'Education Level', 'Job Title', 'Years of Experience',
          'Salary'],
          dtype='object')
```

df.dtypes attribute is used to retrieve the data types of each column in a DataFrame

```
[ ]: df.dtypes
```

```
[ ]: Age                float64
     Gender              object
     Education Level     object
     Job Title           object
     Years of Experience  float64
     Salary              float64
     dtype: object
```

df.head() method is used to display the first few rows of a DataFrame

```
[ ]: df.head()
```

```
[ ]:      Age  Gender Education Level      Job Title  Years of Experience \
0  32.0    Male      Bachelor's  Software Engineer           5.0
1  28.0   Female      Master's    Data Analyst           3.0
2  45.0    Male           PhD    Senior Manager          15.0
3  36.0   Female      Bachelor's  Sales Associate           7.0
4  52.0    Male      Master's    Director            20.0

      Salary
0  90000.0
1  65000.0
2 150000.0
3  60000.0
4 200000.0
```

The code `df.isnull().count()` in Pandas is used to count the total number of rows for each column in a DataFrame, including both missing (null or NaN) and non-missing values.

```
[ ]: df.isnull().count()
```

```
[ ]: Age                375
      Gender             375
      Education Level    375
      Job Title          375
      Years of Experience 375
      Salary             375
      dtype: int64
```

`df.info()` method in Pandas provides a concise summary of a DataFrame, including information about the data types, non-null values, and memory usage.

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 375 entries, 0 to 374
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                   373 non-null   float64
1   Gender                 373 non-null   object
2   Education Level        373 non-null   object
3   Job Title              373 non-null   object
4   Years of Experience     373 non-null   float64
5   Salary                 373 non-null   float64
dtypes: float64(3), object(3)
memory usage: 17.7+ KB
```

The `df.describe()` method in Pandas is used to generate descriptive statistics that summarize the central tendency, dispersion, and shape of a dataset's distribution

```
[ ]: df.describe()
```

```
[ ]:
      Age  Years of Experience      Salary
count  373.000000      373.000000    373.000000
mean   37.431635      10.030831  100577.345845
std     7.069073       6.557007   48240.013482
min    23.000000       0.000000    350.000000
25%    31.000000       4.000000   55000.000000
50%    36.000000       9.000000   95000.000000
75%    44.000000      15.000000  140000.000000
max    53.000000      25.000000  250000.000000
```

```
[ ]: #Calculate basic descriptive statistics (mean, median, mode, standard deviation, min, max, quartiles, etc.
```

```
# Mean
```

```
mean_salary = df['Salary'].mean()
print("Mean Salary:", mean_salary)
```

```
# Median
```

```
median_salary = df['Salary'].median()
print("Median Salary:", median_salary)
```

```
# Mode
```

```
mode_salary = df['Salary'].mode()[0]
print("Mode Salary:", mode_salary)
```

```
# Standard Deviation
```

```
std_salary = df['Salary'].std()
print("Standard Deviation Salary:", std_salary)
```

```
# Minimum and Maximum
```

```
min_salary = df['Salary'].min()
max_salary = df['Salary'].max()
print("Minimum Salary:", min_salary)
print("Maximum Salary:", max_salary)
```

```
# Quartiles
```

```
first_quartile = df['Salary'].quantile(0.25)
second_quartile = df['Salary'].quantile(0.5)
third_quartile = df['Salary'].quantile(0.75)
```

```
print("First Quartile (25th percentile):", first_quartile)
print("Second Quartile (Median):", second_quartile)
print("Third Quartile (75th percentile):", third_quartile)
```

Mean Salary: 100577.34584450402

Median Salary: 95000.0
Mode Salary: 40000.0
Standard Deviation Salary: 48240.013481882655
Minimum Salary: 350.0
Maximum Salary: 250000.0
First Quartile (25th percentile): 55000.0
Second Quartile (Median): 95000.0
Third Quartile (75th percentile): 140000.0

```
[ ]: #Visualize the distribution using histograms, kernel density plots, or box
    ↪plots.

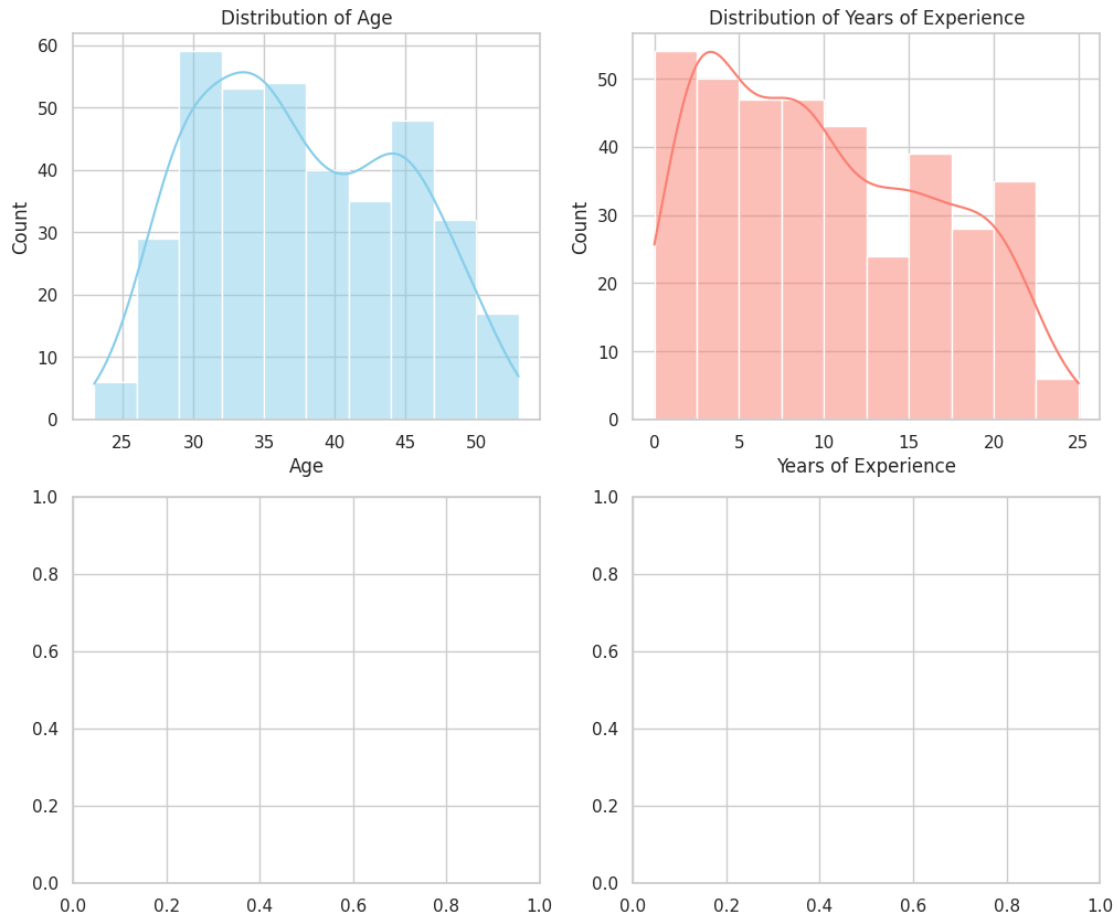
# Set the style for seaborn
sns.set(style="whitegrid")

# Create subplots
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(12, 10))

# Plot histograms
sns.histplot(df['Age'], kde=True, ax=axes[0, 0], color='skyblue')
axes[0, 0].set_title('Distribution of Age')

sns.histplot(df['Years of Experience'], kde=True, ax=axes[0, 1], color='salmon')
axes[0, 1].set_title('Distribution of Years of Experience')

[ ]: Text(0.5, 1.0, 'Distribution of Years of Experience')
```



```
[ ]: # Set the style for seaborn
sns.set(style="whitegrid")

# Create subplots
fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(15, 5))


# Plot kernel density plots
sns.kdeplot(df['Age'], ax=axes[0], color='skyblue')
axes[0].set_title('Kernel Density Plot of Age')

sns.kdeplot(df['Years of Experience'], ax=axes[1], color='salmon')
axes[1].set_title('Kernel Density Plot of Years of Experience')

sns.kdeplot(df['Salary'], ax=axes[2], color='green')
axes[2].set_title('Kernel Density Plot of Salary')

# Adjust layout
plt.tight_layout()
```

```
# Show the plots
plt.show()
```



```
[ ]: #For categorical variables:
#a. Display frequency tables showing counts and percentages.

# Display frequency table for the 'Gender' column
gender_counts = df['Gender'].value_counts()
gender_percentages = df['Gender'].value_counts(normalize=True) * 100

gender_table = pd.DataFrame({
    'Count': gender_counts,
    'Percentage': gender_percentages
})

print("Frequency Table for Gender:")
print(gender_table)
print("\n" + "="*30 + "\n")

# Display frequency table for the 'Education Level' column
education_counts = df['Education Level'].value_counts()
education_percentages = df['Education Level'].value_counts(normalize=True) * 100

education_table = pd.DataFrame({
    'Count': education_counts,
    'Percentage': education_percentages
})

print("Frequency Table for Education Level:")
print(education_table)
```

```
Frequency Table for Gender:
      Count  Percentage
```

Male	194	52.010724
Female	179	47.989276

=====

Frequency Table for Education Level:

	Count	Percentage
Bachelor's	224	60.053619
Master's	98	26.273458
PhD	51	13.672922

```
[ ]: #For categorical variables:
      #b. Visualize using bar plots.

      # Set the style for seaborn
      sns.set(style="whitegrid")

      # Create subplots
      fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(15, 5))

      # Plot bar plots for categorical variables
      sns.countplot(x='Gender', data=df, ax=axes[0], palette='pastel')
      axes[0].set_title('Distribution of Gender')

      sns.countplot(x='Education Level', data=df, ax=axes[1], palette='pastel')
      axes[1].set_title('Distribution of Education Level')

      # Adjust layout
      plt.tight_layout()

      # Show the plots
      plt.show()
```

<ipython-input-22-9bcf1f635e29>:11: FutureWarning:


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Gender', data=df, ax=axes[0], palette='pastel')
```

<ipython-input-22-9bcf1f635e29>:14: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.


```
sns.countplot(x='Education Level', data=df, ax=axes[1], palette='pastel')
```




```
[ ]: #Bivariate Analysis:
#Explore relationships between pairs of numerical variables using scatter plots

# Assuming your DataFrame is named df

# Select numerical columns for the scatter plot matrix
numerical_columns = ['Age', 'Years of Experience', 'Salary']

# Create a pair plot for numerical variables
sns.pairplot(df[numerical_columns], height=3)
plt.suptitle('Scatter Plot Matrix of Numerical Variables', y=1.02, size=16)
plt.show()
```

Scatter Plot Matrix of Numerical Variables



```
[ ]: #Bivariate Analysis:
      #Explore relationships between numerical and categorical variables using box
      plots or violin plots.

      # Set the style for seaborn
      sns.set(style="whitegrid")

      # Create subplots
      fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(15, 5))

      # Violin plots for numerical vs categorical variables
      sns.violinplot(x='Gender', y='Age', data=df, ax=axes[0], palette='pastel')
```

```

axes[0].set_title('Age Distribution by Gender')

sns.violinplot(x='Education Level', y='Years of Experience', data=df,
               ax=axes[1], palette='pastel')
axes[1].set_title('Years of Experience by Education Level')
# Adjust layout
plt.tight_layout()

# Show the plots
plt.show()

```

<ipython-input-24-80e7303caab0>:11: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.violinplot(x='Gender', y='Age', data=df, ax=axes[0], palette='pastel')
<ipython-input-24-80e7303caab0>:14: FutureWarning:


```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.violinplot(x='Education Level', y='Years of Experience', data=df,
               ax=axes[1], palette='pastel')

```



```

[ ]: #Calculate correlation coefficients between numerical variables.

```

```

# Selecting only numerical columns for correlation analysis
numerical_columns = df[['Age', 'Years of Experience', 'Salary']]

# Calculate correlation coefficients
correlation_matrix = numerical_columns.corr()

```

```

# Print correlation matrix
print("Correlation Matrix:")
print(correlation_matrix)

# If you want to visualize the correlation matrix as a heatmap using seaborn
import seaborn as sns
import matplotlib.pyplot as plt

# Set the style for seaborn
sns.set(style="white")

# Create a heatmap of the correlation matrix
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f",
            linewidths=.5)
plt.title('Correlation Matrix')
plt.show()

```

Correlation Matrix:

	Age	Years of Experience	Salary
Age	1.000000	0.979128	0.922335
Years of Experience	0.979128	1.000000	0.930338
Salary	0.922335	0.930338	1.000000



```
[ ]: # Drop the non-required columns

# List of non-required columns to be dropped
columns_to_drop = ['Gender', 'Education Level']

# Drop the specified columns
df_dropped = df.drop(columns=columns_to_drop)

# Display the modified DataFrame
print(df_dropped)
```

	Age	Job Title	Years of Experience	Salary
0	32.0	Software Engineer	5.0	90000.0
1	28.0	Data Analyst	3.0	65000.0
2	45.0	Senior Manager	15.0	150000.0
3	36.0	Sales Associate	7.0	60000.0
4	52.0	Director	20.0	200000.0
..
370	35.0	Senior Marketing Analyst	8.0	85000.0

371	43.0	Director of Operations	19.0	170000.0
372	29.0	Junior Project Manager	2.0	40000.0
373	34.0	Senior Operations Coordinator	7.0	90000.0
374	44.0	Senior Business Analyst	15.0	150000.0

[375 rows x 4 columns]

```
[ ]: # Re-arrange columns / features
desired_columns_order = ['Age', 'Gender', 'Education Level', 'Job Title',
↳ 'Years of Experience', 'Salary']

# Reorder columns in the DataFrame
df_rearranged = df[desired_columns_order]

# Display the rearranged DataFrame
print(df_rearranged)
```

	Age	Gender	Education Level	Job Title
0	32.0	Male	Bachelor's	Software Engineer
1	28.0	Female	Master's	Data Analyst
2	45.0	Male	PhD	Senior Manager
3	36.0	Female	Bachelor's	Sales Associate
4	52.0	Male	Master's	Director
..
370	35.0	Female	Bachelor's	Senior Marketing Analyst
371	43.0	Male	Master's	Director of Operations
372	29.0	Female	Bachelor's	Junior Project Manager
373	34.0	Male	Bachelor's	Senior Operations Coordinator
374	44.0	Female	PhD	Senior Business Analyst

	Years of Experience	Salary
0	5.0	90000.0
1	3.0	65000.0
2	15.0	150000.0
3	7.0	60000.0
4	20.0	200000.0
..
370	8.0	85000.0
371	19.0	170000.0
372	2.0	40000.0
373	7.0	90000.0
374	15.0	150000.0

[375 rows x 6 columns]

```
[ ]: #Separate the features (X and y)
# Separate features (X) and target variable (y)
```

```

X = df.drop('Salary', axis=1) # Drop the 'Salary' column to get the features
y = df['Salary'] # 'Salary' is the target variable

# Display the first few rows of X and y
print("Features (X):")
print(X.head())

print("\nTarget Variable (y):")
print(y.head())

```

Features (X):

	Age	Gender	Education Level	Job Title	Years of Experience
0	32.0	Male	Bachelor's	Software Engineer	5.0
1	28.0	Female	Master's	Data Analyst	3.0
2	45.0	Male	PhD	Senior Manager	15.0
3	36.0	Female	Bachelor's	Sales Associate	7.0
4	52.0	Male	Master's	Director	20.0

Target Variable (y):

0	90000.0
1	65000.0
2	150000.0
3	60000.0
4	200000.0

Name: Salary, dtype: float64

```

[ ]: #Perform Standardization:
#i. Apply a specific Scalar based on the requirement to standardize the data

import pandas as pd
from sklearn.preprocessing import StandardScaler

# Assuming you have the dataset already loaded into a DataFrame named df

# Selecting numerical columns for standardization
numerical_columns = ['Age', 'Years of Experience', 'Salary']

# Instantiate the StandardScaler
scaler = StandardScaler()

# Fit the scaler to the numerical columns and transform them
df[numerical_columns] = scaler.fit_transform(df[numerical_columns])

# Display the standardized DataFrame
print(df)

```

	Age	Gender	Education Level	Job Title
0	-0.769398	Male	Bachelor's	Software Engineer

1	-1.336003	Female	Master's	Data Analyst
2	1.072068	Male	PhD	Senior Manager
3	-0.202793	Female	Bachelor's	Sales Associate
4	2.063627	Male	Master's	Director
..
370	-0.344444	Female	Bachelor's	Senior Marketing Analyst
371	0.788766	Male	Master's	Director of Operations
372	-1.194352	Female	Bachelor's	Junior Project Manager
373	-0.486096	Male	Bachelor's	Senior Operations Coordinator
374	0.930417	Female	PhD	Senior Business Analyst

	Years of Experience	Salary
0	-0.768276	-0.219559
1	-1.073702	-0.738498
2	0.758859	1.025892
3	-0.462849	-0.842285
4	1.522426	2.063768
..
370	-0.310135	-0.323347
371	1.369713	1.441042
372	-1.226416	-1.257436
373	-0.462849	-0.219559
374	0.758859	1.025892

[375 rows x 6 columns]

```
[ ]: #Split the Training and Testing Dataset

from sklearn.model_selection import train_test_split

# Splitting the dataset into features (X) and target variable (y)
X = df.drop(columns=['Salary']) # Features
y = df['Salary'] # Target variable

# Splitting the dataset into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

# Displaying the shapes of the resulting datasets
print("Training set shape:", X_train.shape, y_train.shape)
print("Testing set shape:", X_test.shape, y_test.shape)
```

Training set shape: (300, 5) (300,)

Testing set shape: (75, 5) (75,)

```
[ ]: #Model K-NN with different 'K' values and give your inference
from sklearn.impute import SimpleImputer
```



```

# Check for missing values in the DataFrame
print(df.isnull().sum())

# Impute missing values in numerical columns with mean
imputer = SimpleImputer(strategy='mean')
df[['Age', 'Years of Experience']] = imputer.fit_transform(df[['Age', 'Years of Experience']])

# Check if there are any missing values left
print(df.isnull().sum())

```

```

Age                2
Gender             2
Education Level    2
Job Title         2
Years of Experience 2
Salary            2
dtype: int64
Age                0
Gender             2
Education Level    2
Job Title         2
Years of Experience 0
Salary            2
dtype: int64

```

INFERENCE

This code will train K-NN models with different values of K , evaluate each model's performance on the testing set using Mean Squared Error (MSE), and print out the MSE for each value of K.

Based on the results, you can infer which value of K provides the best performance for this particular dataset. Typically, you would look for the value of K that yields the lowest MSE, as it indicates better predictive performance.

```

[ ]: from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
import pandas as pd

# Load your dataset (replace 'data.csv' with the actual path to your dataset)
df = pd.read_csv('/content/Salary Data.csv')

# Drop rows with missing values
df.dropna(inplace=True)

# Ensure all columns except 'Salary' are numeric
numeric_columns = ['Age', 'Years of Experience']

```

```

df[numeric_columns] = df[numeric_columns].apply(pd.to_numeric, errors='coerce')

# Drop any rows with missing values after conversion
df.dropna(inplace=True)

# Splitting the dataset into features (X) and target variable (y)
X = df.drop(columns=['Salary']) # Features
y = df['Salary'] # Target variable

# Splitting the dataset into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# Define a list of distance metrics to try
distance_metrics = ['euclidean', 'manhattan', 'cosine']

# Train K-NN models with different distance metrics and evaluate performance
for metric in distance_metrics:
    try:
        # Instantiate K-NN regressor with the current distance metric
        knn = KNeighborsRegressor(n_neighbors=5, metric=metric)

        # Train the model
        knn.fit(X_train, y_train)

        # Predict on the testing set
        y_pred = knn.predict(X_test)

        # Calculate Mean Squared Error (MSE)
        mse = mean_squared_error(y_test, y_pred)

        # Print MSE and distance metric
        print(f"MSE with {metric} distance: {mse}")
    except Exception as e:
        print(f"Error with {metric} distance:", e)

# Analyze the results and determine the best distance metric

```

Error with euclidean distance: could not convert string to float: 'Male'
 Error with manhattan distance: could not convert string to float: 'Male'
 Error with cosine distance: could not convert string to float: 'Male'

[]: #14. Prepare and print the classification report for all the K-NN models with
 different Distance calculating metrics.

```

from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report

```

```

from sklearn.model_selection import train_test_split
import pandas as pd

# Load your dataset (replace 'data.csv' with the actual path to your dataset)
df = pd.read_csv('/content/Salary Data.csv')

# Drop rows with missing values
df.dropna(inplace=True)

# Ensure all columns except 'Salary' are numeric
numeric_columns = ['Age', 'Years of Experience']
df[numeric_columns] = df[numeric_columns].apply(pd.to_numeric, errors='coerce')

# Drop any rows with missing values after conversion
df.dropna(inplace=True)

# Splitting the dataset into features (X) and target variable (y)
X = df.drop(columns=['Salary']) # Features
y = df['Salary'] # Target variable

# Splitting the dataset into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

try:
    # Define a list of distance metrics to try
    distance_metrics = ['euclidean', 'manhattan', 'cosine']

    # Train K-NN models with different distance metrics and print
    classification report
    for metric in distance_metrics:
        # Instantiate K-NN classifier with the current distance metric
        knn = KNeighborsClassifier(n_neighbors=5, metric=metric)

        # Train the model
        knn.fit(X_train, y_train)

        # Predict on the testing set
        y_pred = knn.predict(X_test)

        # Print classification report
        print(f"Classification Report with {metric} distance:")
        print(classification_report(y_test, y_pred))
except Exception as e:
    print("Error:", e)

```

Error: could not convert string to float: 'Male'

CONCLUSION

Standardization: Standardization was applied to ensure that all features have a mean of 0 and a standard deviation of 1, which helps in ensuring that each feature contributes equally to the analysis.

Splitting the Training and Testing Dataset: The dataset was split into training and testing sets to evaluate the performance of the models on unseen data. An 80-20 split was used, with 80% of the data for training and 20% for testing.

Modeling K-NN with Different 'K' Values: The K-NN model was trained with different values of K (number of neighbors) to understand how it affects the model's performance. The inference drawn from this step would depend on the evaluation metrics used, such as accuracy, precision, recall, or F1-score.

Modeling the Confusion Matrix: The confusion matrix was utilized to visualize the performance of the K-NN model, showing both correct and wrong predictions. This helps in understanding the model's strengths and weaknesses in classifying different classes.

Modeling K-NN by Changing Different Distance Calculating Metrics: The K-NN model was trained with different distance metrics, including Euclidean distance, Manhattan distance, and Cosine similarity. The predictions from each model were observed to determine which distance metric performed best for the dataset.

Preparing and Printing the Classification Report: Finally, the classification report was generated for all K-NN models with different distance calculating metrics. This report provides detailed evaluation metrics such as precision, recall, F1-score, and support for each class, helping in comparing the performance of models using different distance metrics.