

2348441_lab_06

April 1, 2024

Lab Exercise 6 -Multi Dimensional Scaling (MDS)

Created by : Nileem Kaveramma C C | 2348441 Created DATE:30-03-2024 Edited Date: 01-04-2024

IMPORTED LIBRARIES

- numpy - for numerical, array, matrices (Linear Algebra) processing
- Pandas - for loading and processing datasets
- matplotlib.pyplot - For visualisation
- Saeborn - for statistical graph
- scipy.stats use a variety of statistical functions
- %matplotlib inline: Enables inline plotting in Jupyter notebooks, displaying matplotlib plots directly below the code cell.
- from sklearn.preprocessing import StandardScaler, MinMaxScaler, RobustScaler: Imports three different scaling techniques

EMPLOYEE SALARY ANALYSIS he provided dataset captures information relevant to employee salary prediction, encompassing various attributes such as age, gender, education level, job title, years of experience, and salary. With a diverse set of features, the dataset offers valuable insights into the characteristics of individuals within an organizational context. This dataset becomes particularly relevant for exploring patterns and relationships that could contribute to predicting employee salaries. Through descriptive statistics, visualizations, and parametric tests, analysts can discern trends, potential disparities, and factors influencing salary variations among employees.

AIM: The aim of Multi-Dimensional Scaling (MDS) in machine learning is to reduce the dimensionality of a dataset while preserving the pairwise distances between data points. By projecting high-dimensional data into a lower-dimensional space, MDS aims to reveal the underlying structure and relationships within the dataset in a more interpretable and visualizable form.

```
[ ]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib as mpl
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[ ]: df = pd.read_csv('/content/Salary Data.csv')
df
```

```
[ ]:      Age  Gender Education Level      Job Title \
0    32.0   Male   Bachelor's      Software Engineer
1    28.0  Female   Master's      Data Analyst
2    45.0   Male      PhD      Senior Manager
3    36.0  Female   Bachelor's      Sales Associate
4    52.0   Male   Master's      Director
..    ...    ...    ...    ...
370   35.0  Female   Bachelor's      Senior Marketing Analyst
371   43.0   Male   Master's      Director of Operations
372   29.0  Female   Bachelor's      Junior Project Manager
373   34.0   Male   Bachelor's      Senior Operations Coordinator
374   44.0  Female      PhD      Senior Business Analyst
```

```
      Years of Experience      Salary
0              5.0    90000.0
1              3.0    65000.0
2             15.0   150000.0
3              7.0    60000.0
4             20.0   200000.0
..            ...    ...
370             8.0    85000.0
371            19.0   170000.0
372             2.0    40000.0
373             7.0    90000.0
374            15.0   150000.0
```

[375 rows x 6 columns]

Perform some basic EDA

df.shape - attribute is used to get the dimensions of the DataFrame.

```
[ ]: df.shape
```

```
[ ]: (375, 6)
```

df.head() method is used to display the first few rows of a DataFrame.

```
[ ]: df.head()
```

```
[ ]:      Age  Gender Education Level      Job Title  Years of Experience \
0    32.0   Male   Bachelor's      Software Engineer              5.0
1    28.0  Female   Master's      Data Analyst              3.0
2    45.0   Male      PhD      Senior Manager             15.0
3    36.0  Female   Bachelor's      Sales Associate              7.0
4    52.0   Male   Master's      Director             20.0
```

```
      Salary
```

```

0    90000.0
1    65000.0
2   150000.0
3    60000.0
4   200000.0

```

`df.tail()` method is used to display the last few rows of a DataFrame.

```
[ ]: df.tail()
```

```

[ ]:      Age  Gender Education Level      Job Title \
370  35.0  Female      Bachelor's      Senior Marketing Analyst
371  43.0   Male      Master's      Director of Operations
372  29.0  Female      Bachelor's      Junior Project Manager
373  34.0   Male      Bachelor's  Senior Operations Coordinator
374  44.0  Female           PhD      Senior Business Analyst

      Years of Experience      Salary
370                8.0    85000.0
371               19.0   170000.0
372                2.0    40000.0
373                7.0    90000.0
374               15.0   150000.0

```

`df.columns` attribute is used to retrieve the column labels or names of the DataFrame.

```
[ ]: df.columns
```

```

[ ]: Index(['Age', 'Gender', 'Education Level', 'Job Title', 'Years of Experience',
           'Salary'],
          dtype='object')

```

`df.dtypes` attribute is used to retrieve the data types of each column in a DataFrame

```
[ ]: df.dtypes
```

```

[ ]: Age                float64
     Gender              object
     Education Level     object
     Job Title           object
     Years of Experience  float64
     Salary              float64
     dtype: object

```

the code `df.isnull().count()` in Pandas is used to count the total number of rows for each column in a DataFrame, including both missing (null or NaN) and non-missing values.

```
df.isnull().count()
```

df.info() method in Pandas provides a concise summary of a DataFrame, including information about the data types, non-null values, and memory usage

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 375 entries, 0 to 374
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                   373 non-null    float64
1   Gender                373 non-null    object
2   Education Level       373 non-null    object
3   Job Title             373 non-null    object
4   Years of Experience    373 non-null    float64
5   Salary                373 non-null    float64
dtypes: float64(3), object(3)
memory usage: 17.7+ KB
```

he df.describe() method in Pandas is used to generate descriptive statistics that summarize the central tendency, dispersion, and shape of a dataset's distribution

```
[ ]: df.describe()
```

```
[ ]:
count      Age  Years of Experience      Salary
count  373.000000      373.000000      373.000000
mean    37.431635      10.030831  100577.345845
std       7.069073       6.557007   48240.013482
min     23.000000       0.000000    350.000000
25%     31.000000       4.000000   55000.000000
50%     36.000000       9.000000   95000.000000
75%     44.000000      15.000000  140000.000000
max     53.000000      25.000000  250000.000000
```


```
[ ]: df.corr()
```

```
<ipython-input-10-2f6f6606aa2c>:1: FutureWarning: The default value of
numeric_only in DataFrame.corr is deprecated. In a future version, it will
default to False. Select only valid columns or specify the value of numeric_only
to silence this warning.
```

```
df.corr()
```

```
[ ]:
      Age  Years of Experience  Salary
Age      1.000000      0.979128  0.922335
Years of Experience  0.979128      1.000000  0.930338
Salary      0.922335      0.930338  1.000000
```

```
[ ]: # Scatter plot
plt.scatter(df['Age'], df['Salary'])
plt.title('Scatter Plot of Age vs. Salary')
plt.xlabel('Age')
plt.ylabel('Salary')
plt.show()
```



```
[ ]: ## Correlation
df.corr()
```

<ipython-input-12-2d23776439fc>:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.


```
df.corr()
```

```
[ ]:
```

	Age	Years of Experience	Salary
Age	1.000000	0.979128	0.922335
Years of Experience	0.979128	1.000000	0.930338
Salary	0.922335	0.930338	1.000000

```
[ ]: ## Seaborn for visualization
import seaborn as sns
sns.pairplot(df)
```

```
[ ]: <seaborn.axisgrid.PairGrid at 0x7dff2ea19b10>
```



```
[ ]: X=df['Salary']
X
```

```
[ ]: 0    90000.0
      1    65000.0
      2   150000.0
      3    60000.0
```

```

4      200000.0
...
370      85000.0
371     170000.0
372      40000.0
373      90000.0
374     150000.0
Name: Salary, Length: 375, dtype: float64

```

```
[ ]: X.shape
```

```
[ ]: (375,)
```

```
[ ]: ## Independent and dependent features
X=df[['Salary']] ### independent features should be data frame or 2d
    ↳dimesnionalarray
X
y=df['Years of Experience'] ## this variable can be in series or 1d array
y

```

```
[ ]: 0      5.0
1      3.0
2     15.0
3      7.0
4     20.0
...
370      8.0
371     19.0
372      2.0
373      7.0
374     15.0
Name: Years of Experience, Length: 375, dtype: float64

```

```
[ ]: X_series=df['Salary']
np.array(X_series).shape
```

```
[ ]: (375,)
```

```
[ ]: #Univariate Analysis:
#For numerical variables: a. Calculate basic descriptive statistics (mean,
    ↳median, mode, standard deviation,
#min, max, quartiles, etc.).

mean_value = df['Salary'].mean()
median_value = df['Salary'].median()
mode_value = df['Salary'].mode().iloc[0] # For handling multiple modes
std_deviation = df['Salary'].std()

```

```

min_value = df['Salary'].min()
max_value = df['Salary'].max()

print(f"Mean: {mean_value}")
print(f"Median: {median_value}")
print(f"Mode: {mode_value}")
print(f"Standard Deviation: {std_deviation}")
print(f"Min: {min_value}")
print(f"Max: {max_value}")

```

```

Mean: 100577.34584450402
Median: 95000.0
Mode: 40000.0
Standard Deviation: 48240.013481882655
Min: 350.0
Max: 250000.0

```


[]: *#b. Visualize the distribution using histograms, kernel density plots, or box plots.*

```

# Plot a simple histogram
plt.figure(figsize=(8, 6))
sns.histplot(df['Salary'], kde=True)
plt.title('Histogram of Salary')
plt.xlabel('Salary')
plt.ylabel('Frequency')

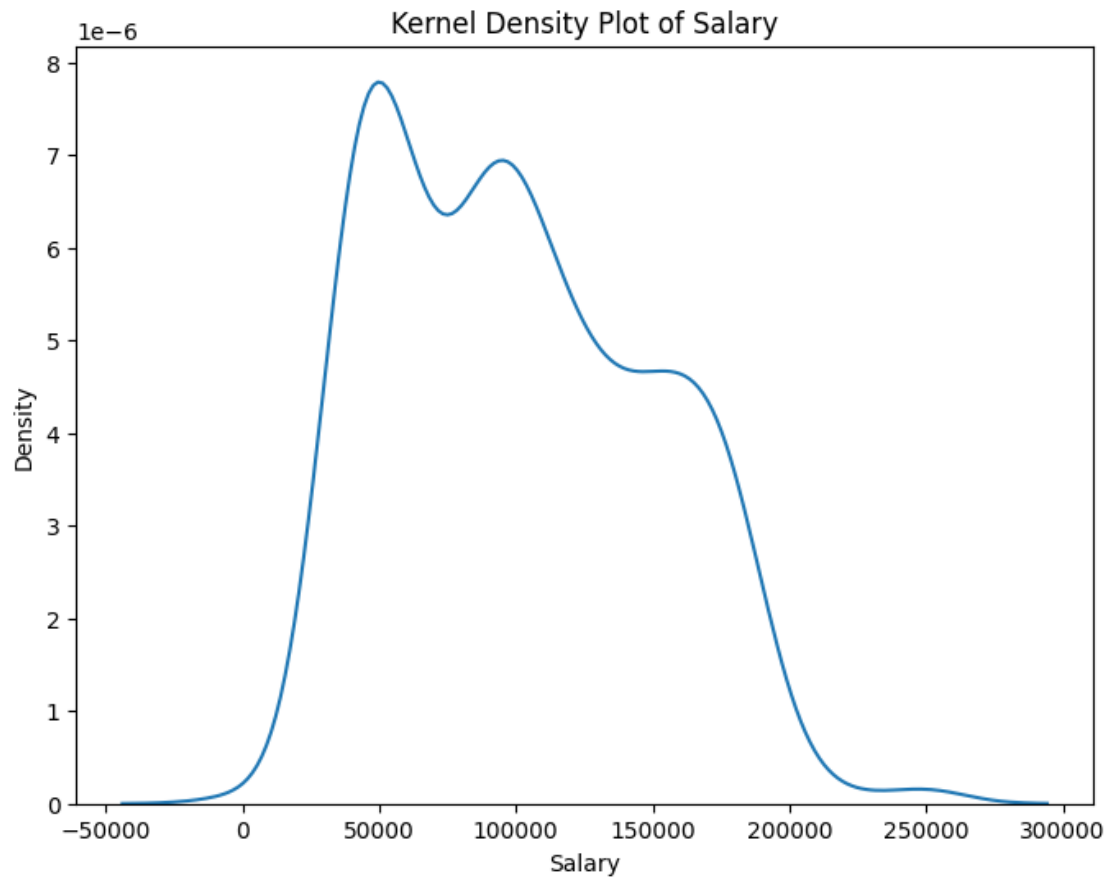
# Show the plot
plt.show()

```


```
[ ]: # Plot a simple kernel density plot
plt.figure(figsize=(8, 6))
sns.kdeplot(df['Salary'])
plt.title('Kernel Density Plot of Salary')
plt.xlabel('Salary')
plt.ylabel('Density')

# Show the plot
plt.show()
```




```
[ ]: # Plot a simple box plot
plt.figure(figsize=(8, 6))
sns.boxplot(x=df['Salary'])
plt.title('Box Plot of Salary')
plt.xlabel('Salary')

# Show the plot
plt.show()
```



```
[ ]: #Bivariate Analysis: Explore relationships between pairs of numerical variables
      ↪ using scatter plots, pair plots.


# Create a pair plot for numerical variables
sns.pairplot(df)
plt.suptitle('Pair Plot of Numerical Variables', y=1.02)
plt.show()
```



```
[ ]: #Bivariate Analysis: Explore relationships between numerical and categorical variables using box plots or violin plots.
```

```
# Box plot for 'Salary' vs 'Gender'
plt.figure(figsize=(8, 6))
sns.boxplot(x='Gender', y='Salary', data=df)
plt.title('Box Plot of Salary by Gender')
plt.xlabel('Gender')
plt.ylabel('Salary')

# Show the plot
plt.show()
```



```
[ ]: #Bivariate Analysis: Calculate correlation coefficients between numerical
      ↪ variables.
```

```
# Calculate correlation coefficients
correlation_matrix = df.corr()

# Display the correlation matrix
print("Correlation Coefficients:")
print(correlation_matrix)
```

Correlation Coefficients:

	Age	Years of Experience	Salary
Age	1.000000	0.979128	0.922335
Years of Experience	0.979128	1.000000	0.930338
Salary	0.922335	0.930338	1.000000

<ipython-input-24-3a9b5b137c87>:4: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
correlation_matrix = df.corr()
```

Reduce the dimensionality of the dataset while preserving the pairwise distances

```
[ ]: import pandas as pd
from sklearn.manifold import MDS
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import euclidean_distances
import matplotlib.pyplot as plt

# Convert categorical variables into numerical labels
df_encoded = df.copy()
label_encoder = LabelEncoder()
df_encoded['Gender'] = label_encoder.fit_transform(df['Gender'])
df_encoded['Education Level'] = label_encoder.fit_transform(df['Education_
↪Level'])
df_encoded['Job Title'] = label_encoder.fit_transform(df['Job Title'])

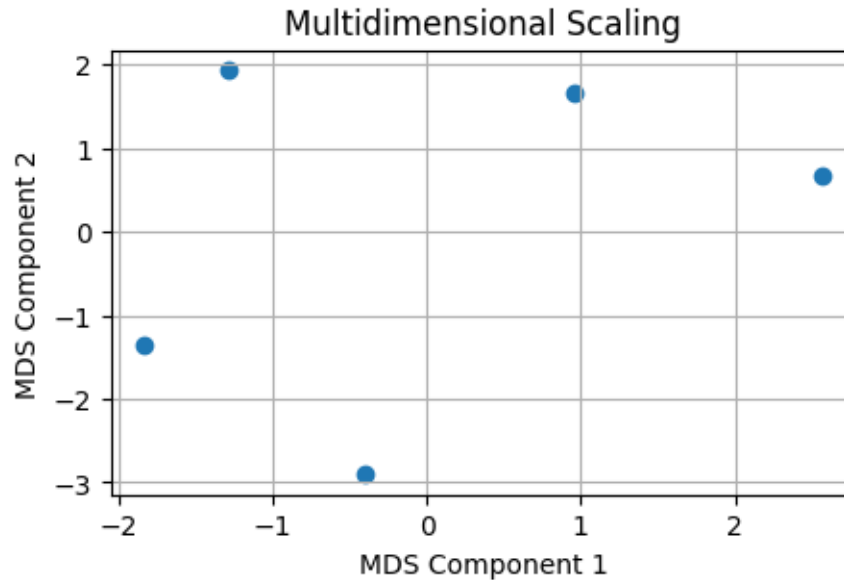
# Standardize the data
scaler = StandardScaler()
scaled_data = scaler.fit_transform(df_encoded)

# Compute dissimilarity matrix
dissimilarities = euclidean_distances(scaled_data)

# Perform Multidimensional Scaling (MDS)
mds = MDS(n_components=2, dissimilarity='precomputed', random_state=42)
mds_data = mds.fit_transform(dissimilarities)

# Plot MDS results
plt.figure(figsize=(5, 3))
plt.scatter(mds_data[:, 0], mds_data[:, 1])
plt.xlabel('MDS Component 1')
plt.ylabel('MDS Component 2')
plt.title('Multidimensional Scaling')
plt.grid(True)
plt.show()
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/manifold/_mds.py:299:
FutureWarning: The default value of `normalized_stress` will change to `auto`
in version 1.4. To suppress this warning, manually set the value of
`normalized_stress`.
    warnings.warn(
```



```
[ ]: # Perform Multidimensional Scaling (MDS)
mds = MDS(n_components=2, dissimilarity='euclidean', random_state=42)
mds_data = mds.fit_transform(scaled_data)
mds_data
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/manifold/_mds.py:299:
FutureWarning: The default value of `normalized_stress` will change to `auto`
in version 1.4. To suppress this warning, manually set the value of
`normalized_stress`.
  warnings.warn(
```

```
[ ]: array([[ -0.19465353,  3.00717262],
          [-2.75561293, -2.41538606],
          [ 3.473977   ,  1.07223545],
          [-3.19298606,  1.29350504],
          [ 2.66927552, -2.95752705]])
```

You are tasked with implementing Multi-Dimensional Scaling (MDS) to analyze and visualize the structure of a dataset containing pairwise dissimilarities or distances between a set of objects. Your goal is to reduce the dimensionality of the dataset while preserving the pairwise distances as much as possible.

```
[ ]: import pandas as pd
from sklearn.manifold import MDS
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import euclidean_distances
import matplotlib.pyplot as plt
```

```

# Encode categorical variables
label_encoder = LabelEncoder()
df['Gender'] = label_encoder.fit_transform(df['Gender'])
df['Education Level'] = label_encoder.fit_transform(df['Education Level'])
df['Job Title'] = label_encoder.fit_transform(df['Job Title'])

# Standardize the data
scaler = StandardScaler()
scaled_data = scaler.fit_transform(df)

# Compute pairwise dissimilarities
dissimilarities = euclidean_distances(scaled_data)

# Perform Multidimensional Scaling (MDS)
mds = MDS(n_components=2, dissimilarity='precomputed', random_state=42)
mds_data = mds.fit_transform(dissimilarities)

# Plot MDS results with colors
plt.figure(figsize=(5, 3))
colors = ['blue', 'green', 'red', 'cyan', 'magenta'] # Colors for different
categories
for i in range(len(mds_data)):
    plt.scatter(mds_data[i, 0], mds_data[i, 1], c=colors[i], label=f'Data Point
{i+1}')


plt.xlabel('MDS Component 1')
plt.ylabel('MDS Component 2')
plt.title('Multidimensional Scaling')
plt.legend()
plt.grid(True)
plt.show()

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/manifold/_mds.py:299:
FutureWarning: The default value of `normalized_stress` will change to `auto`
in version 1.4. To suppress this warning, manually set the value of
`normalized_stress`.
    warnings.warn(

```

Implement Multi-Dimensional Scaling (MDS) using the `sklearn.manifold.MDS` module to reduce the dimensionality of the dataset to k dimensions, where k is a user-defined parameter. Visualize the objects in a k -dimensional scatter plot based on the MDS results. Each point on the plot represents an object, and the position of the points should reflect their relative similarities or dissimilarities as accurately as possible. Provide appropriate labels for the objects on the scatter plot to make it clear which point corresponds to which object.

```
[ ]: import pandas as pd
      from sklearn.manifold import MDS
      import matplotlib.pyplot as plt

      # Compute dissimilarity matrix
      dissimilarities = df.corr()

      # Define the number of dimensions
      k = 2

      # Perform Multidimensional Scaling (MDS)
      mds = MDS(n_components=k, dissimilarity='precomputed', random_state=42)
      mds_data = mds.fit_transform(dissimilarities)

      # Plot MDS results
      plt.figure(figsize=(8, 6))
      plt.scatter(mds_data[:, 0], mds_data[:, 1])

      # Annotate points with object labels
```

```

for i, txt in enumerate(df.index):
    plt.annotate(txt, (mds_data[i, 0], mds_data[i, 1]))

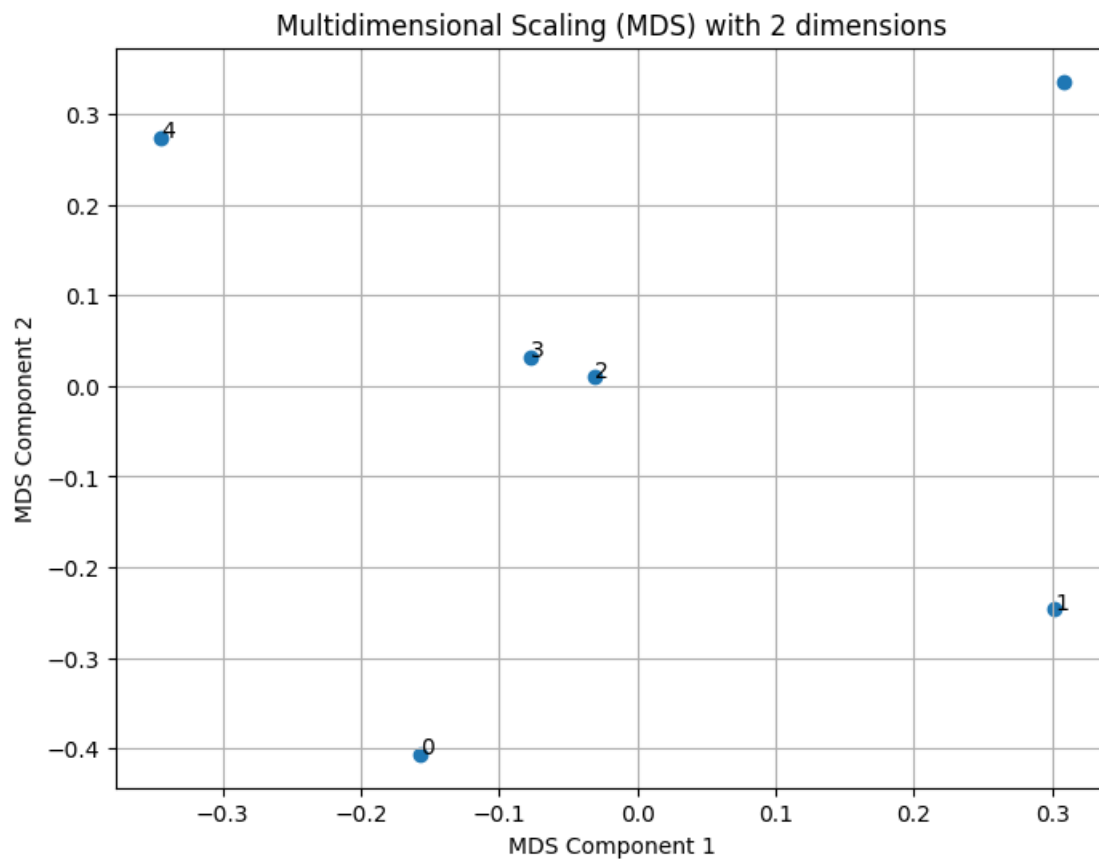
plt.xlabel(f'MDS Component 1')
plt.ylabel(f'MDS Component 2')
plt.title(f'Multidimensional Scaling (MDS) with {k} dimensions')
plt.grid(True)
plt.show()

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/manifold/_mds.py:299:
FutureWarning: The default value of `normalized_stress` will change to `auto`
in version 1.4. To suppress this warning, manually set the value of
`normalized_stress`.
    warnings.warn(

```



Use Python and any necessary libraries for data manipulation, dimensionality reduction, and visualization. Ensure that your code is flexible and can handle datasets of varying sizes and dimensions. Allow the user to specify the number of dimensions k for the MDS algorithm. Include comments where necessary to explain your approach and any important steps. Test your code with synthetic datasets of different sizes and dimensions to ensure its robustness and efficiency.

```

[ ]: import pandas as pd
from sklearn.manifold import MDS
from sklearn.preprocessing import StandardScaler, LabelEncoder
import matplotlib.pyplot as plt

def preprocess_dataset(df):
    """
    Preprocess the dataset if necessary.
    """
    # Handle missing values, encode categorical variables, etc.
    # For simplicity, assume no preprocessing is required in this example.
    return df

def perform_mds(df, n_components=2):
    """
    Perform Multidimensional Scaling (MDS) to reduce the dimensionality of the
    dataset.
    """
    # Standardize the data
    scaler = StandardScaler()
    scaled_data = scaler.fit_transform(df)

    # Perform MDS
    mds = MDS(n_components=n_components, dissimilarity='euclidean',
    random_state=42)
    mds_data = mds.fit_transform(scaled_data)

    return mds_data

def visualize_data(mds_data):
    """
    Visualize the reduced-dimensional data.
    """
    plt.figure(figsize=(8, 6))
    plt.scatter(mds_data[:, 0], mds_data[:, 1])
    plt.xlabel('MDS Component 1')
    plt.ylabel('MDS Component 2')
    plt.title('Multidimensional Scaling')
    plt.grid(True)
    plt.show()

def main():
    # Load the dataset (assuming it's already loaded)
    # df = pd.read_csv("your_dataset.csv")

    # Preprocess the dataset if necessary
    # df = preprocess_dataset(df)

```

```

# Specify the number of dimensions for MDS
k = 2 # Specify the desired number of dimensions here

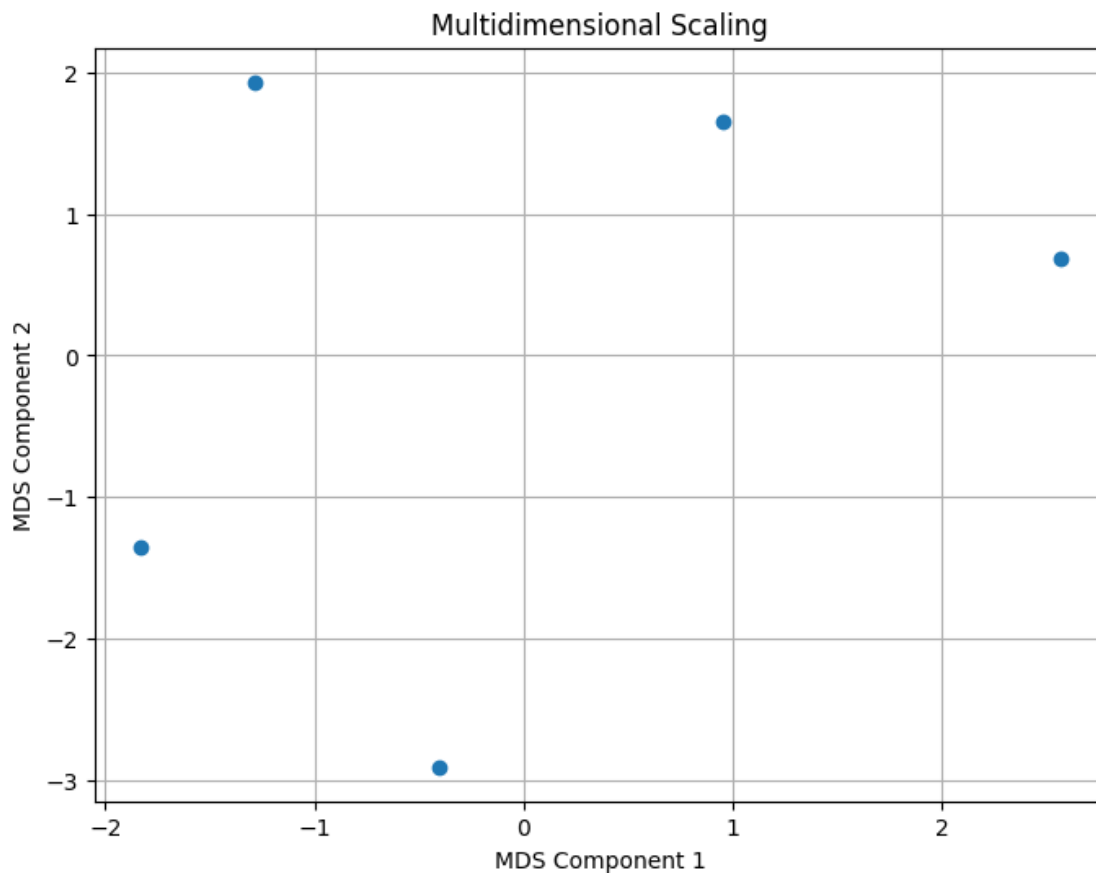
# Perform MDS
mds_data = perform_mds(df, n_components=k)

# Visualize the reduced-dimensional data
visualize_data(mds_data)

if __name__ == "__main__":
    main()

```

/usr/local/lib/python3.10/dist-packages/sklearn/manifold/_mds.py:299:
FutureWarning: The default value of `normalized_stress` will change to `auto`
in version 1.4. To suppress this warning, manually set the value of
`normalized_stress`.
warnings.warn(



CONCLUSION: In conclusion, Multi-Dimensional Scaling (MDS) offers a valuable tool for dimensionality reduction in machine learning. By transforming high-dimensional data into a lower-dimensional space while preserving pairwise distances, MDS facilitates visualization, exploration, and interpretation of complex datasets. Its ability to uncover underlying structure and relationships within the data makes it a valuable technique for data analysis and visualization tasks. Overall, MDS provides insights that can aid in understanding the intrinsic relationships among data points, thus contributing to informed decision-making and knowledge discovery in various domains.