

# 2348441\_lab\_4

March 9, 2024

## Lab Exercise 4 -Regression Analysis

Created by : Nileem Kaveramma C C | 2348441 Created DATE:08-03-2024 Edited Date: 09-03-2024

### IMPORTED LIBRARIES

- numpy - for numerical, array, matrices (Linear Algebra) processing
- Pandas - for loading and processing datasets
- matplotlib.pyplot - For visualisation
- Saeborn - for statistical graph
- scipy.stats use a variety of statistical functions
- %matplotlib inline: Enables inline plotting in Jupyter notebooks, displaying matplotlib plots directly below the code cell.
- from sklearn.preprocessing import StandardScaler, MinMaxScaler, RobustScaler: Imports three different scaling techniques
- from sklearn.decomposition import PCA: Imports the Principal Component Analysis (PCA) module from scikit-learn for dimensionality reduction and feature extraction.
- from sklearn.preprocessing import StandardScaler: Imports StandardScaler from scikit-learn, a method for standardizing numerical features by removing the mean and scaling to unit variance.

EMPLOYEE SALARY ANALYSIS he provided dataset captures information relevant to employee salary prediction, encompassing various attributes such as age, gender, education level, job title, years of experience, and salary. With a diverse set of features, the dataset offers valuable insights into the characteristics of individuals within an organizational context. This dataset becomes particularly relevant for exploring patterns and relationships that could contribute to predicting employee salaries. Through descriptive statistics, visualizations, and parametric tests, analysts can discern trends, potential disparities, and factors influencing salary variations among employees.

AIM : The aim of Principal Component Analysis (PCA) is to reduce the dimensionality of a dataset while retaining as much of the original variability as possible. By transforming the original features into a new set of uncorrelated variables called principal components, PCA simplifies data representation, aids in identifying patterns, and facilitates efficient analysis and visualization, particularly useful in high-dimensional datasets.

### PROCEDURE:

Data Collection: Gather a dataset with numerical features for which dimensionality reduction is desired.

Data Standardization: Standardize the data by subtracting the mean and scaling to unit variance. This step ensures that all features contribute equally to the analysis.

**Covariance Matrix Computation:** Calculate the covariance matrix of the standardized data. The covariance matrix represents the relationships between different features.

**Eigenvalue and Eigenvector Calculation:** Find the eigenvalues and corresponding eigenvectors of the covariance matrix. These eigenvectors represent the principal components, and the eigenvalues indicate their importance.

**Sort Eigenvalues and Eigenvectors:** Arrange the eigenvalues and their corresponding eigenvectors in descending order. The principal components with higher eigenvalues capture more variance in the data.

**Selection of Principal Components:** Determine the number of principal components to retain based on the explained variance. This decision often involves selecting components that explain a significant percentage (e.g., 95%) of the total variance.

**Projection of Data onto Principal Components:** Project the original data onto the selected principal components to obtain a reduced-dimensional representation of the dataset.

**Analysis and Visualization:** Analyze the transformed data using the selected principal components. Visualize the results to gain insights into the structure and patterns within the dataset.

**Interpretation of Principal Components:** Interpret the principal components in the context of the original features to understand the underlying patterns and relationships in the data.

**Validation and Iteration:** Validate the results and iterate the procedure as needed, adjusting the number of retained components or exploring alternative preprocessing steps.

1) Import all the necessary libraries

```
[ ]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib as mpl
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
from sklearn.preprocessing import StandardScaler, MinMaxScaler, RobustScaler
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
```

2)Loading the dataset

```
[ ]: df = pd.read_csv('/content/Salary Data.csv')
df
```

```
[ ]:      Age  Gender Education Level      Job Title \
0    32.0   Male   Bachelor's      Software Engineer
1    28.0  Female   Master's      Data Analyst
2    45.0   Male      PhD      Senior Manager
3    36.0  Female   Bachelor's      Sales Associate
4    52.0   Male   Master's      Director
..    ...    ...      ...      ...
```

370	35.0	Female	Bachelor's	Senior Marketing Analyst
371	43.0	Male	Master's	Director of Operations
372	29.0	Female	Bachelor's	Junior Project Manager
373	34.0	Male	Bachelor's	Senior Operations Coordinator
374	44.0	Female	PhD	Senior Business Analyst

	Years of Experience	Salary
0	5.0	90000.0
1	3.0	65000.0
2	15.0	150000.0
3	7.0	60000.0
4	20.0	200000.0
..	...	...
370	8.0	85000.0
371	19.0	170000.0
372	2.0	40000.0
373	7.0	90000.0
374	15.0	150000.0

[375 rows x 6 columns]

3). Perform some basic EDA

```
[ ]: columns_names=df.columns.tolist()
print("Columns names:")
print(columns_names)
```

Columns names:

```
['Age', 'Gender', 'Education Level', 'Job Title', 'Years of Experience',
'Salary']
```

df.columns.tolist() fetches all the columns and then convert it into list type.This step is just to check out all the column names in our data.Columns are also called as features of our datasets.

df.shape - attribute is used to get the dimensions of the DataFrame.

```
[ ]: df.shape
```

```
[ ]: (375, 6)
```

df.head() method is used to display the first few rows of a DataFrame.

```
[ ]: df.head()
```

```
[ ]:      Age  Gender Education Level      Job Title  Years of Experience  \
0   32.0   Male      Bachelor's  Software Engineer           5.0
1   28.0  Female      Master's    Data Analyst           3.0
2   45.0   Male          PhD    Senior Manager          15.0
3   36.0  Female      Bachelor's  Sales Associate           7.0
```

4	52.0	Male	Master's	Director	20.0
---	------	------	----------	----------	------

	Salary
0	90000.0
1	65000.0
2	150000.0
3	60000.0
4	200000.0

df.tail() method is used to display the last few rows of a DataFrame.

```
[ ]: df.tail()
```

```
[ ]:      Age  Gender Education Level      Job Title \
370  35.0  Female   Bachelor's      Senior Marketing Analyst
371  43.0   Male    Master's      Director of Operations
372  29.0  Female   Bachelor's      Junior Project Manager
373  34.0   Male   Bachelor's  Senior Operations Coordinator
374  44.0  Female      PhD      Senior Business Analyst

      Years of Experience      Salary
370                8.0    85000.0
371               19.0   170000.0
372                2.0    40000.0
373                7.0    90000.0
374               15.0   150000.0
```

df.columns attribute is used to retrieve the column labels or names of the DataFrame.

```
[ ]: df.columns
```

```
[ ]: Index(['Age', 'Gender', 'Education Level', 'Job Title', 'Years of Experience',
           'Salary'],
          dtype='object')
```

df.dtypes attribute is used to retrieve the data types of each column in a DataFrame

```
[ ]: df.dtypes
```

```
[ ]: Age                float64
     Gender              object
     Education Level     object
     Job Title           object
     Years of Experience  float64
     Salary              float64
     dtype: object
```

the code df.isnull().count() in Pandas is used to count the total number of rows for each column in a DataFrame, including both missing (null or NaN) and non-missing values.

```
[ ]: df.isnull().count()
```

```
[ ]: Age                375
      Gender            375
      Education Level   375
      Job Title         375
      Years of Experience 375
      Salary            375
      dtype: int64
```

df.info() method in Pandas provides a concise summary of a DataFrame, including information about the data types, non-null values, and memory usage

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 375 entries, 0 to 374
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                   373 non-null   float64
1   Gender                373 non-null   object
2   Education Level       373 non-null   object
3   Job Title             373 non-null   object
4   Years of Experience    373 non-null   float64
5   Salary                373 non-null   float64
dtypes: float64(3), object(3)
memory usage: 17.7+ KB
```

The df.describe() method in Pandas is used to generate descriptive statistics that summarize the central tendency, dispersion, and shape of a dataset's distribution

```
[ ]: df.describe()
```

```
[ ]:
count    Age  Years of Experience  Salary
mean    37.431635    10.030831  100577.345845
std      7.069073     6.557007   48240.013482
min     23.000000     0.000000    350.000000
25%     31.000000     4.000000   55000.000000
50%     36.000000     9.000000   95000.000000
75%     44.000000    15.000000  140000.000000
max     53.000000    25.000000  250000.000000
```

```
[ ]: df.corr()
```

```
<ipython-input-12-2f6f6606aa2c>:1: FutureWarning: The default value of
numeric_only in DataFrame.corr is deprecated. In a future version, it will
default to False. Select only valid columns or specify the value of numeric_only
```


```
to silence this warning.  
df.corr()
```

```
[ ]:      Age  Years of Experience  Salary  
Age      1.000000      0.979128  0.922335  
Years of Experience  0.979128      1.000000  0.930338  
Salary      0.922335      0.930338  1.000000
```

```
[ ]: correlation = df.corr()  
plt.figure(figsize=(10,10))  
sns.heatmap(correlation, vmax=1, square=True,annot=True,cmap='cubehelix')  
  
plt.title('Correlation between different fearures')
```

```
<ipython-input-13-effb445d3340>:1: FutureWarning: The default value of  
numeric_only in DataFrame.corr is deprecated. In a future version, it will  
default to False. Select only valid columns or specify the value of numeric_only  
to silence this warning.  
correlation = df.corr()
```

```
[ ]: Text(0.5, 1.0, 'Correlation between different fearures')
```



Performing some visualisation before moving onto PCA

```
[ ]: grouped_data = df.groupby('Job Title').sum()
      print(grouped_data)
```

Job Title	Age	Years of Experience	Salary
Account Manager	32.0	5.0	75000.0
Accountant	31.0	4.0	55000.0
Administrative Assistant	75.0	18.0	100000.0
Business Analyst	67.0	12.0	155000.0
Business Development Manager	34.0	8.0	90000.0
...	...	...	...

UX Designer	34.0	5.0	80000.0
UX Researcher	27.0	2.0	65000.0
VP of Finance	47.0	19.0	200000.0
VP of Operations	47.0	19.0	190000.0
Web Developer	33.0	6.0	65000.0

[174 rows x 3 columns]

<ipython-input-14-9d5f552b0293>:1: FutureWarning: The default value of numeric\_only in DataFrameGroupBy.sum is deprecated. In a future version, numeric\_only will default to False. Either specify numeric\_only or select only columns which should be valid for the function.

```
grouped_data = df.groupby('Job Title').sum()
```

```
[ ]: df['Job Title'].unique()
```

```
[ ]: array(['Software Engineer', 'Data Analyst', 'Senior Manager',
'Sales Associate', 'Director', 'Marketing Analyst',
'Product Manager', 'Sales Manager', 'Marketing Coordinator',
'Senior Scientist', 'Software Developer', 'HR Manager',
'Financial Analyst', 'Project Manager', 'Customer Service Rep',
'Operations Manager', 'Marketing Manager', 'Senior Engineer',
'Data Entry Clerk', 'Sales Director', 'Business Analyst',
'VP of Operations', 'IT Support', 'Recruiter', 'Financial Manager',
'Social Media Specialist', 'Software Manager', 'Junior Developer',
'Senior Consultant', 'Product Designer', 'CEO', 'Accountant',
'Data Scientist', 'Marketing Specialist', 'Technical Writer',
'HR Generalist', 'Project Engineer', 'Customer Success Rep',
'Sales Executive', 'UX Designer', 'Operations Director',
'Network Engineer', 'Administrative Assistant',
'Strategy Consultant', 'Copywriter', 'Account Manager',
'Director of Marketing', 'Help Desk Analyst',
'Customer Service Manager', 'Business Intelligence Analyst',
'Event Coordinator', 'VP of Finance', 'Graphic Designer',
'UX Researcher', 'Social Media Manager', 'Director of Operations',
'Senior Data Scientist', 'Junior Accountant',
'Digital Marketing Manager', 'IT Manager',
'Customer Service Representative', 'Business Development Manager',
'Senior Financial Analyst', 'Web Developer', 'Research Director',
'Technical Support Specialist', 'Creative Director',
'Senior Software Engineer', 'Human Resources Director',
'Content Marketing Manager', 'Technical Recruiter',
'Sales Representative', 'Chief Technology Officer',
'Junior Designer', 'Financial Advisor', 'Junior Account Manager',
'Senior Project Manager', 'Principal Scientist',
'Supply Chain Manager', 'Senior Marketing Manager',
'Training Specialist', 'Research Scientist',
'Junior Software Developer', 'Public Relations Manager',
```



```

'Operations Analyst', 'Product Marketing Manager',
'Senior HR Manager', 'Junior Web Developer',
'Senior Project Coordinator', 'Chief Data Officer',
'Digital Content Producer', 'IT Support Specialist',
'Senior Marketing Analyst', 'Customer Success Manager',
'Senior Graphic Designer', 'Software Project Manager',
'Supply Chain Analyst', 'Senior Business Analyst',
'Junior Marketing Analyst', 'Office Manager', 'Principal Engineer',
'Junior HR Generalist', 'Senior Product Manager',
'Junior Operations Analyst', 'Senior HR Generalist',
'Sales Operations Manager', 'Senior Software Developer',
'Junior Web Designer', 'Senior Training Specialist',
'Senior Research Scientist', 'Junior Sales Representative',
'Junior Marketing Manager', 'Junior Data Analyst',
'Senior Product Marketing Manager', 'Junior Business Analyst',
'Senior Sales Manager', 'Junior Marketing Specialist',
'Junior Project Manager', 'Senior Accountant', 'Director of Sales',
'Junior Recruiter', 'Senior Business Development Manager',
'Senior Product Designer', 'Junior Customer Support Specialist',
'Senior IT Support Specialist', 'Junior Financial Analyst',
'Senior Operations Manager', 'Director of Human Resources',
'Junior Software Engineer', 'Senior Sales Representative',
'Director of Product Management', 'Junior Copywriter',
'Senior Marketing Coordinator', 'Senior Human Resources Manager',
'Junior Business Development Associate', 'Senior Account Manager',
'Senior Researcher', 'Junior HR Coordinator',
'Director of Finance', 'Junior Marketing Coordinator', nan,
'Junior Data Scientist', 'Senior Operations Analyst',
'Senior Human Resources Coordinator', 'Senior UX Designer',
'Junior Product Manager', 'Senior Marketing Specialist',
'Senior IT Project Manager', 'Senior Quality Assurance Analyst',
'Director of Sales and Marketing', 'Senior Account Executive',
'Director of Business Development', 'Junior Social Media Manager',
'Senior Human Resources Specialist', 'Senior Data Analyst',
'Director of Human Capital', 'Junior Advertising Coordinator',
'Junior UX Designer', 'Senior Marketing Director',
'Senior IT Consultant', 'Senior Financial Advisor',
'Junior Business Operations Analyst',
'Junior Social Media Specialist',
'Senior Product Development Manager', 'Junior Operations Manager',
'Senior Software Architect', 'Junior Research Scientist',
'Senior Financial Manager', 'Senior HR Specialist',
'Senior Data Engineer', 'Junior Operations Coordinator',
'Director of HR', 'Senior Operations Coordinator',
'Junior Financial Advisor', 'Director of Engineering'],
dtype=object)

```

```
[ ]: groupby_sales=df.groupby('Job Title').mean()
groupby_sales
```

<ipython-input-16-455fffa59049>:1: FutureWarning: The default value of numeric\_only in DataFrameGroupBy.mean is deprecated. In a future version, numeric\_only will default to False. Either specify numeric\_only or select only columns which should be valid for the function.

```
groupby_sales=df.groupby('Job Title').mean()
```

```
[ ]:
      Age  Years of Experience  Salary
Job Title
Account Manager      32.0         5.0  75000.0
Accountant           31.0         4.0  55000.0
Administrative Assistant 37.5         9.0  50000.0
Business Analyst     33.5         6.0  77500.0
Business Development Manager 34.0         8.0  90000.0
...
UX Designer          34.0         5.0  80000.0
UX Researcher        27.0         2.0  65000.0
VP of Finance        47.0        19.0 200000.0
VP of Operations     47.0        19.0 190000.0
Web Developer        33.0         6.0  65000.0
```

[174 rows x 3 columns]

```
[ ]: #Calculate basic descriptive statistics (mean, median, mode, standard deviation, min, max, quartiles, etc.
```

```
# Mean
```

```
mean_salary = df['Salary'].mean()
print("Mean Salary:", mean_salary)
```

```
# Median
```

```
median_salary = df['Salary'].median()
print("Median Salary:", median_salary)
```

```
# Mode
```

```
mode_salary = df['Salary'].mode()[0]
print("Mode Salary:", mode_salary)
```

```
# Standard Deviation
```

```
std_salary = df['Salary'].std()
print("Standard Deviation Salary:", std_salary)
```

```
# Minimum and Maximum
```

```
min_salary = df['Salary'].min()
max_salary = df['Salary'].max()
```

```

print("Minimum Salary:", min_salary)
print("Maximum Salary:", max_salary)

# Quartiles
first_quartile = df['Salary'].quantile(0.25)
second_quartile = df['Salary'].quantile(0.5)
third_quartile = df['Salary'].quantile(0.75)

print("First Quartile (25th percentile):", first_quartile)
print("Second Quartile (Median):", second_quartile)
print("Third Quartile (75th percentile):", third_quartile)

```

```

Mean Salary: 100577.34584450402
Median Salary: 95000.0
Mode Salary: 40000.0
Standard Deviation Salary: 48240.013481882655
Minimum Salary: 350.0
Maximum Salary: 250000.0
First Quartile (25th percentile): 55000.0
Second Quartile (Median): 95000.0
Third Quartile (75th percentile): 140000.0

```

```

[ ]: #Visualize the distribution using histograms, kernel density plots, or box
    ↪ plots.

# Set the style for seaborn
sns.set(style="whitegrid")

# Create subplots
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(12, 10))

# Plot histograms
sns.histplot(df['Age'], kde=True, ax=axes[0, 0], color='skyblue')
axes[0, 0].set_title('Distribution of Age')

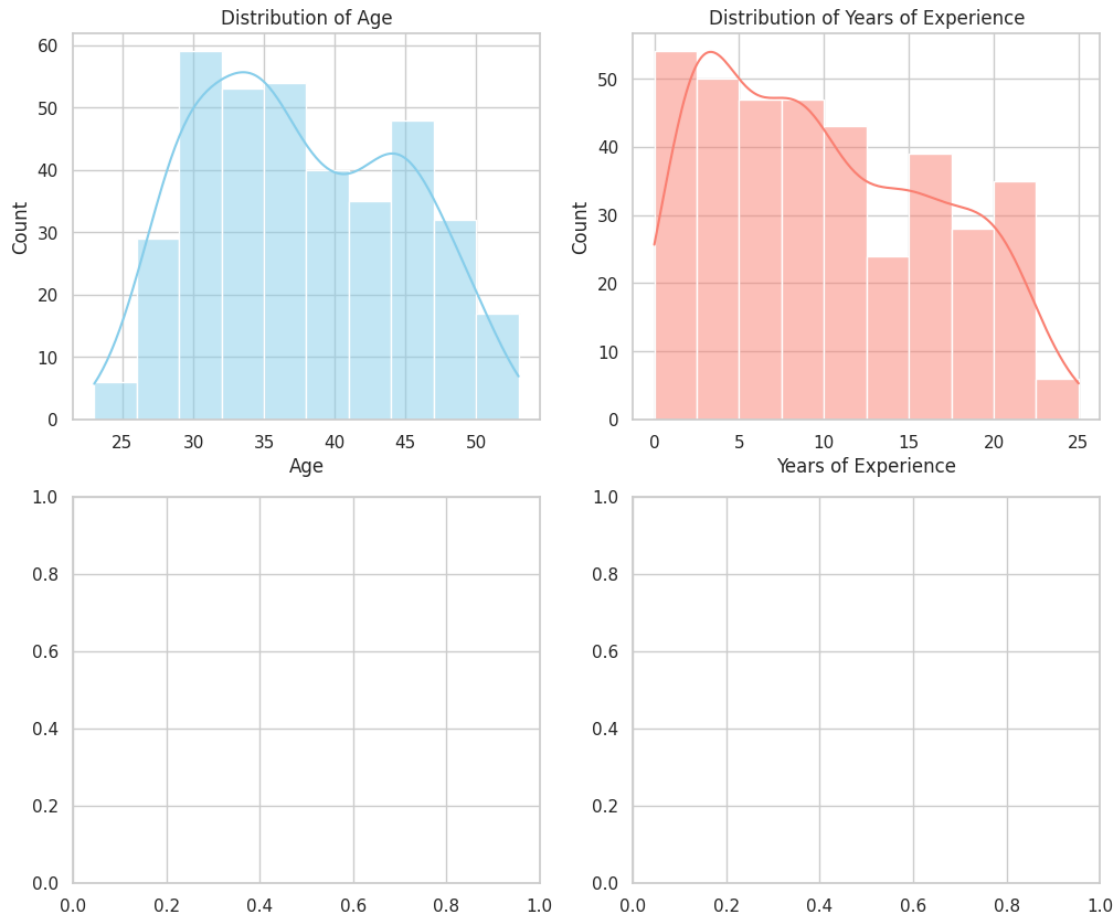
sns.histplot(df['Years of Experience'], kde=True, ax=axes[0, 1], color='salmon')
axes[0, 1].set_title('Distribution of Years of Experience')

```

```

[ ]: Text(0.5, 1.0, 'Distribution of Years of Experience')

```



```
[ ]:
```

```
[ ]: # Set the style for seaborn
sns.set(style="whitegrid")

# Create subplots
fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(15, 5))


# Plot kernel density plots
sns.kdeplot(df['Age'], ax=axes[0], color='skyblue')
axes[0].set_title('Kernel Density Plot of Age')

sns.kdeplot(df['Years of Experience'], ax=axes[1], color='salmon')
axes[1].set_title('Kernel Density Plot of Years of Experience')

sns.kdeplot(df['Salary'], ax=axes[2], color='green')
axes[2].set_title('Kernel Density Plot of Salary')
```

```
# Adjust layout
plt.tight_layout()

# Show the plots
plt.show()
```



```
[ ]: #For categorical variables:
#a. Display frequency tables showing counts and percentages.

# Display frequency table for the 'Gender' column
gender_counts = df['Gender'].value_counts()
gender_percentages = df['Gender'].value_counts(normalize=True) * 100

gender_table = pd.DataFrame({
    'Count': gender_counts,
    'Percentage': gender_percentages
})

print("Frequency Table for Gender:")
print(gender_table)
print("\n" + "="*30 + "\n")

# Display frequency table for the 'Education Level' column
education_counts = df['Education Level'].value_counts()
education_percentages = df['Education Level'].value_counts(normalize=True) * 100

education_table = pd.DataFrame({
    'Count': education_counts,
    'Percentage': education_percentages
})

print("Frequency Table for Education Level:")
print(education_table)
```

Frequency Table for Gender:

	Count	Percentage
Male	194	52.010724
Female	179	47.989276

=====

Frequency Table for Education Level:

	Count	Percentage
Bachelor's	224	60.053619
Master's	98	26.273458
PhD	51	13.672922

```
[ ]: #For categorical variables:
#b. Visualize using bar plots.

# Set the style for seaborn
sns.set(style="whitegrid")

# Create subplots
fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(15, 5))

# Plot bar plots for categorical variables
sns.countplot(x='Gender', data=df, ax=axes[0], palette='pastel')
axes[0].set_title('Distribution of Gender')

sns.countplot(x='Education Level', data=df, ax=axes[1], palette='pastel')
axes[1].set_title('Distribution of Education Level')


# Adjust layout
plt.tight_layout()

# Show the plots
plt.show()
```

<ipython-input-21-9bcf1f635e29>:11: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.


```
sns.countplot(x='Gender', data=df, ax=axes[0], palette='pastel')
```

<ipython-input-21-9bcf1f635e29>:14: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same


effect.

```
sns.countplot(x='Education Level', data=df, ax=axes[1], palette='pastel')
```



```
[ ]: #Bivariate Analysis:  
#Explore relationships between pairs of numerical variables using scatter plots  
  
# Assuming your DataFrame is named df  
  
# Select numerical columns for the scatter plot matrix  
numerical_columns = ['Age', 'Years of Experience', 'Salary']  
  
# Create a pair plot for numerical variables  
sns.pairplot(df[numerical_columns], height=3)  
plt.suptitle('Scatter Plot Matrix of Numerical Variables', y=1.02, size=16)  
plt.show()
```

Scatter Plot Matrix of Numerical Variables



```
[ ]: #Bivariate Analysis:
#Explore relationships between numerical and categorical variables using box
plots or violin plots.

# Set the style for seaborn
sns.set(style="whitegrid")

# Create subplots
fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(15, 5))

# Violin plots for numerical vs categorical variables
sns.violinplot(x='Gender', y='Age', data=df, ax=axes[0], palette='pastel')
```



```

axes[0].set_title('Age Distribution by Gender')

sns.violinplot(x='Education Level', y='Years of Experience', data=df,
               ax=axes[1], palette='pastel')
axes[1].set_title('Years of Experience by Education Level')
# Adjust layout
plt.tight_layout()

# Show the plots
plt.show()

```

<ipython-input-23-80e7303caab0>:11: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.violinplot(x='Gender', y='Age', data=df, ax=axes[0], palette='pastel')
<ipython-input-23-80e7303caab0>:14: FutureWarning:


```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.violinplot(x='Education Level', y='Years of Experience', data=df,
               ax=axes[1], palette='pastel')

```



```

[ ]: #Calculate correlation coefficients between numerical variables.

```

```

# Selecting only numerical columns for correlation analysis
numerical_columns = df[['Age', 'Years of Experience', 'Salary']]

# Calculate correlation coefficients
correlation_matrix = numerical_columns.corr()

```

```

# Print correlation matrix
print("Correlation Matrix:")
print(correlation_matrix)

# If you want to visualize the correlation matrix as a heatmap using seaborn
import seaborn as sns
import matplotlib.pyplot as plt

# Set the style for seaborn
sns.set(style="white")

# Create a heatmap of the correlation matrix
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f",
            linewidths=.5)
plt.title('Correlation Matrix')
plt.show()

```

Correlation Matrix:

	Age	Years of Experience	Salary
Age	1.000000	0.979128	0.922335
Years of Experience	0.979128	1.000000	0.930338
Salary	0.922335	0.930338	1.000000



```
[ ]: # Drop the non-required columns

# List of non-required columns to be dropped
columns_to_drop = ['Gender', 'Education Level']

# Drop the specified columns
df_dropped = df.drop(columns=columns_to_drop)

# Display the modified DataFrame
print(df_dropped)
```

	Age	Job Title	Years of Experience	Salary
0	32.0	Software Engineer	5.0	90000.0
1	28.0	Data Analyst	3.0	65000.0
2	45.0	Senior Manager	15.0	150000.0
3	36.0	Sales Associate	7.0	60000.0
4	52.0	Director	20.0	200000.0
...	...	...	...	...
370	35.0	Senior Marketing Analyst	8.0	85000.0

371	43.0	Director of Operations	19.0	170000.0
372	29.0	Junior Project Manager	2.0	40000.0
373	34.0	Senior Operations Coordinator	7.0	90000.0
374	44.0	Senior Business Analyst	15.0	150000.0

[375 rows x 4 columns]

```
[ ]: # Re-arrange columns / features
desired_columns_order = ['Age', 'Gender', 'Education Level', 'Job Title',
↳ 'Years of Experience', 'Salary']

# Reorder columns in the DataFrame
df_rearranged = df[desired_columns_order]

# Display the rearranged DataFrame
print(df_rearranged)
```

	Age	Gender	Education Level	Job Title
0	32.0	Male	Bachelor's	Software Engineer
1	28.0	Female	Master's	Data Analyst
2	45.0	Male	PhD	Senior Manager
3	36.0	Female	Bachelor's	Sales Associate
4	52.0	Male	Master's	Director
..	...	...	...	...
370	35.0	Female	Bachelor's	Senior Marketing Analyst
371	43.0	Male	Master's	Director of Operations
372	29.0	Female	Bachelor's	Junior Project Manager
373	34.0	Male	Bachelor's	Senior Operations Coordinator
374	44.0	Female	PhD	Senior Business Analyst

	Years of Experience	Salary
0	5.0	90000.0
1	3.0	65000.0
2	15.0	150000.0
3	7.0	60000.0
4	20.0	200000.0
..	...	...
370	8.0	85000.0
371	19.0	170000.0
372	2.0	40000.0
373	7.0	90000.0
374	15.0	150000.0

[375 rows x 6 columns]

```
[ ]: #Separate the features (X and y)
# Separate features (X) and target variable (y)
```

```

X = df.drop('Salary', axis=1) # Drop the 'Salary' column to get the features
y = df['Salary'] # 'Salary' is the target variable

# Display the first few rows of X and y
print("Features (X):")
print(X.head())

print("\nTarget Variable (y):")
print(y.head())

```

Features (X):

	Age	Gender	Education Level	Job Title	Years of Experience
0	32.0	Male	Bachelor's	Software Engineer	5.0
1	28.0	Female	Master's	Data Analyst	3.0
2	45.0	Male	PhD	Senior Manager	15.0
3	36.0	Female	Bachelor's	Sales Associate	7.0
4	52.0	Male	Master's	Director	20.0

Target Variable (y):

```

0    90000.0
1    65000.0
2   150000.0
3    60000.0
4   200000.0

```

Name: Salary, dtype: float64

```

[ ]: #Perform Standardization:Apply Standard Scalar / MinMax Scalar / Robust Scalar
      ↳based on the
      #requirement to standardize the data

# Extract numerical columns for standardization
numerical_columns = ['Age', 'Years of Experience', 'Salary']

# Create a DataFrame containing only numerical columns
numerical_df = df[numerical_columns]

# Initialize the scalers
standard_scaler = StandardScaler()
minmax_scaler = MinMaxScaler()
robust_scaler = RobustScaler()

# Standardize the data using each scaler
standardized_data_standard = standard_scaler.fit_transform(numerical_df)
standardized_data_minmax = minmax_scaler.fit_transform(numerical_df)
standardized_data_robust = robust_scaler.fit_transform(numerical_df)

# Convert the standardized data back to a DataFrame

```

```

standardized_df_standard = pd.DataFrame(standardized_data_standard,
    ↪columns=numerical_columns)
standardized_df_minmax = pd.DataFrame(standardized_data_minmax,
    ↪columns=numerical_columns)
standardized_df_robust = pd.DataFrame(standardized_data_robust,
    ↪columns=numerical_columns)

# Display the standardized DataFrames
print("Standard Scaler:")
print(standardized_df_standard.head())

print("\nMinMax Scaler:")
print(standardized_df_minmax.head())

print("\nRobust Scaler:")
print(standardized_df_robust.head())

```

Standard Scaler:

	Age	Years of Experience	Salary
0	-0.769398	-0.768276	-0.219559
1	-1.336003	-1.073702	-0.738498
2	1.072068	0.758859	1.025892
3	-0.202793	-0.462849	-0.842285
4	2.063627	1.522426	2.063768

MinMax Scaler:

	Age	Years of Experience	Salary
0	0.300000	0.20	0.359103
1	0.166667	0.12	0.258963
2	0.733333	0.60	0.599439
3	0.433333	0.28	0.238935
4	0.966667	0.80	0.799720

Robust Scaler:

	Age	Years of Experience	Salary
0	-0.307692	-0.363636	-0.058824
1	-0.615385	-0.545455	-0.352941
2	0.692308	0.545455	0.647059
3	0.000000	-0.181818	-0.411765
4	1.230769	1.000000	1.235294

**Conventional-way of PCA:**

```

[ ]: #Compute Eigenvectors and Eigenvalues

# Drop rows with missing values
df_cleaned = df.dropna()

```

```

# Extract numerical columns for PCA
numerical_columns = ['Age', 'Years of Experience', 'Salary']

# Create a DataFrame containing only numerical columns
numerical_df = df_cleaned[numerical_columns]

# Standardize the data (optional but often recommended before PCA)
standardized_data = (numerical_df - numerical_df.mean()) / numerical_df.std()

# Check for NaN or Inf values after standardization
if np.any(np.isnan(standardized_data)) or np.any(np.isinf(standardized_data)):
    raise ValueError("NaN or Inf values found in the standardized data. Handle_
    ↪missing values before PCA.")

# Compute the covariance matrix
covariance_matrix = np.cov(standardized_data, rowvar=False)

# Compute eigenvectors and eigenvalues
eigenvalues, eigenvectors = np.linalg.eig(covariance_matrix)

# Display the results
print("Eigenvalues:")
print(eigenvalues)

print("\nEigenvectors:")
print(eigenvectors)

```

Eigenvalues:

```
[2.88809241 0.02056164 0.09134595]
```

Eigenvectors:

```
[[-0.58016321 -0.67969213 -0.4488087 ]
 [-0.58175427  0.73145357 -0.35572129]
 [-0.57006369 -0.05471997  0.81977626]]
```

```

[ ]: #Check the Correlation between features


# Calculate the correlation matrix
correlation_matrix = df.corr()

# Create a heatmap to visualize the correlation matrix
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f',
    ↪linewidths=.5)
plt.title('Correlation Matrix')
plt.show()

```

<ipython-input-38-631b40c7a995>:4: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
correlation_matrix = df.corr()
```



```
[ ]: #Select Principal components with Covariance
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
import pandas as pd

# Assuming your DataFrame is named df

# Handle missing values (fill with the mean of each column)
df_filled = df.fillna(df.mean())

# Selecting numerical columns for PCA
```



```

numerical_columns = df_filled.select_dtypes(include=['float64']).columns

# Extract numerical data
data = df_filled[numerical_columns]

# Standardize the data
scaler = StandardScaler()
scaled_data = scaler.fit_transform(data)

# Apply PCA with covariance
pca = PCA()
principal_components = pca.fit_transform(scaled_data)

# Create a DataFrame with principal components
columns = [f'PC{i+1}' for i in range(principal_components.shape[1])]
principal_df = pd.DataFrame(data=principal_components, columns=columns)

# Display the explained variance ratio for each principal component
print("Explained Variance Ratio:")
print(pca.explained_variance_ratio_)

# Optional: You can also print the cumulative explained variance
print("\nCumulative Explained Variance:")
print(pca.explained_variance_ratio_.cumsum())

# Optional: Print the loadings (coefficients) of each variable on each
# principal component
loadings_df = pd.DataFrame(data=pca.components_.T, index=data.columns,
                           columns=columns)
print("\nLoadings:")
print(loadings_df)

```

Explained Variance Ratio:  
[0.96269747 0.03044865 0.00685388]

Cumulative Explained Variance:  
[0.96269747 0.99314612 1.            ]

Loadings:

	PC1	PC2	PC3
Age	0.580163	-0.448809	-0.679692
Years of Experience	0.581754	-0.355721	0.731454
Salary	0.570064	0.819776	-0.054720

<ipython-input-48-436792dddc3>:9: FutureWarning: The default value of numeric\_only in DataFrame.mean is deprecated. In a future version, it will default to False. In addition, specifying 'numeric\_only=None' is deprecated. Select only valid columns or specify the value of numeric\_only to silence this

```
warning.  
df_filled = df.fillna(df.mean())
```

```
[ ]: #Projection Matrix  
import numpy as np  
  
# Assuming your DataFrame is named df  
  
# Extract numerical columns for simplicity  
numerical_data = df[['Age', 'Years of Experience', 'Salary']].values  
  
# Handle NaN values by replacing them with the mean of each column  
numerical_data = np.nan_to_num(numerical_data, nan=np.nanmean(numerical_data,  
↪axis=0))  
  
# Calculate the covariance matrix  
covariance_matrix = np.cov(numerical_data, rowvar=False)  
  
# Perform eigendecomposition of the covariance matrix  
eigenvalues, eigenvectors = np.linalg.eig(covariance_matrix)  
  
# Sort eigenvectors based on eigenvalues  
sorted_indices = np.argsort(eigenvalues)[::-1]  
eigenvectors_sorted = eigenvectors[:, sorted_indices]  
  
# Choose the number of dimensions for projection (e.g., 2 for 2D projection)  
num_dimensions = 2  
  
# Create the projection matrix  
projection_matrix = eigenvectors_sorted[:, :num_dimensions]  
  
# Project the data onto the subspace defined by the projection matrix  
projected_data = np.dot(numerical_data, projection_matrix)  
  
# Display the projection matrix  
print("Projection Matrix:")  
print(projection_matrix)
```

```
Projection Matrix:  
[[ 1.35158649e-04  7.57625050e-01]  
 [ 1.26455831e-04  6.52690010e-01]  
 [ 9.99999983e-01 -1.84936039e-04]]
```

```
[ ]: #Calculate PCA with Scikit-learn  
  
# Extracting numerical columns for PCA  
numerical_columns = ['Age', 'Years of Experience', 'Salary']
```

```

data_for_pca = df[numerical_columns]

# Handle missing values by imputing with mean (you can choose a different
↳strategy)
imputer = SimpleImputer(strategy='mean')
data_for_pca_imputed = pd.DataFrame(imputer.fit_transform(data_for_pca),
↳columns=numerical_columns)

# Standardize the data
scaler = StandardScaler()
scaled_data = scaler.fit_transform(data_for_pca_imputed)

# Apply PCA
pca = PCA()
pca_result = pca.fit_transform(scaled_data)


# Create a DataFrame with the PCA results
pca_df = pd.DataFrame(data=pca_result, columns=['PC1', 'PC2', 'PC3'])

# Print explained variance ratios
explained_variance_ratios = pca.explained_variance_ratio_
print("Explained Variance Ratios:")
print(explained_variance_ratios)

# Plot the cumulative explained variance
cumulative_variance = explained_variance_ratios.cumsum()
plt.plot(range(1, len(cumulative_variance) + 1), cumulative_variance,
↳marker='o', linestyle='--')
plt.title('Cumulative Explained Variance')
plt.xlabel('Number of Principal Components')
plt.ylabel('Cumulative Explained Variance')
plt.show()

```

Explained Variance Ratios:  
[0.96269747 0.03044865 0.00685388]



**Interpret the results with your conclusion to substantiate that the principal components shall provide equivalent understanding of all the features.**

Explained Variance Ratios:

Higher values indicate more substantial contributions of corresponding principal components to the total variance, reflecting their importance in capturing information from the original features. Cumulative Explained Variance Plot:

Reveals diminishing returns as additional principal components are considered. Common practice is to retain components until a point on the x-axis where further inclusion doesn't significantly enhance cumulative explained variance. This chosen point represents a trade-off between dimensionality reduction and retaining a high percentage of total variance.

### **Conclusion:**

If a small number of principal components can explain a high percentage of the total variance, it suggests that these components provide a condensed representation of the original features. This can be valuable for dimensionality reduction, simplifying the analysis while retaining most of the information.

Look at the cumulative explained variance plot to determine how many principal components are needed to achieve a satisfactory level of data representation. If a small subset of components can

capture a significant portion of the variance, it supports the idea that these components offer an equivalent understanding of the original features.

PCA is particularly useful when dealing with high-dimensional data, and it allows you to focus on the most informative components while discarding less significant ones. The interpretation of these components in the context of your specific data domain is crucial for extracting meaningful insights.